

Activity__ Course 5 Automatidata project lab

May 10, 2023

1 Automatidata project

Course 5 - Regression Analysis: Simplify complex data relationships

The data consulting firm Automatidata has recently hired you as the newest member of their data analytics team. Their newest client, the NYC Taxi and Limousine Commission (New York City TLC), wants the Automatidata team to build a multiple linear regression model for ride durations based on a variety of variables. The team is getting closer to completing the project, having completed an initial plan of action, initial Python coding work, EDA, and A/B testing.

The Automatidata team has reviewed the results of the A/B testing. Now it's time to work on predicting the taxi/ride share trip durations. You've impressed your Automatidata colleagues with your hard work and attention to detail. The data team believes that you are ready to build the regression model and update the client New York City TLC about your progress.

2 Course 5 End-of-course project: Build a multiple linear regression model

In this activity, you will build a multiple linear regression model

The purpose of this project is to demonstrate knowledge of EDA and a multiple linear regression model

The goal is to build a multiple linear regression model and evaluate the model *This activity has three parts:*

Part 1: EDA & Checking Model Assumptions * What are some purposes of EDA before constructing a multiple linear regression model?

Part 2: Model Building and evaluation * What resources do you find yourself using as you complete this stage?

Part 3: Interpreting Model Results

- What key insights emerged from your model(s)?
- What business recommendations do you propose based on the models built?

Recall that you have a helpful tool at your disposal! Refer to the [PACE strategy document](#) to help apply your learning, apply new problem-solving skills, and guide your approach to this project.

3 Build a multiple linear regression model

4 PACE stages

- [Plan] (#scrollTo=psz51YkZVwtN&line=3&uniquifier=1)
- [Analyze] (#scrollTo=mA7Mz_SnI8km&line=4&uniquifier=1)
- [Construct] (#scrollTo=Lca9c8XON8lc&line=2&uniquifier=1)
- [Execute] (#scrollTo=401PgchTPr4E&line=2&uniquifier=1)

5 Pace: Plan

5.1 Task 1. Imports and loading

Import the packages that you've learned are needed for building linear regression models.

```
[89]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from datetime import datetime
from datetime import date
from datetime import timedelta

from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
import sklearn.metrics as metrics
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, r2_score, mean_squared_error
lr=LinearRegression()
```

Pandas is used to load the NYC TLC dataset. As shown in this cell, the dataset has been automatically loaded in for you. You do not need to download the .csv file, or provide more code, in order to access the dataset and proceed with this lab. Please continue with this activity by completing the following instructions.

```
[90]: # Read in data from NYC TLC dataset provided and load into dataframe
df=pd.read_csv("2017_Yellow_Taxi_Trip_Data.csv") # index_col parameter_
↳specified to avoid "Unnamed: 0" column when reading in data from csv
print('Data loaded')
```

Data loaded

6 PACE: Analyze

In this stage, consider the following question where applicable to complete your code response:

- What are some purposes of EDA before constructing a multiple linear regression model?

Improve the performance and accuracy of the model while reducing the possibility of bias in the distribution of the dataset.

6.1 Task 2a. Explore data with EDA

Analyze and discover data, looking for correlations, missing data, outliers, and duplicates.

Start with `.shape` and `.info()`.

```
[91]: print(df.shape)
      print(df.info())
```

```
(22699, 18)
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 22699 entries, 0 to 22698
Data columns (total 18 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Unnamed: 0            22699 non-null  int64
 1   VendorID              22699 non-null  int64
 2   tpep_pickup_datetime  22699 non-null  object
 3   tpep_dropoff_datetime 22699 non-null  object
 4   passenger_count       22699 non-null  int64
 5   trip_distance         22699 non-null  float64
 6   RatecodeID            22699 non-null  int64
 7   store_and_fwd_flag    22699 non-null  object
 8   PULocationID          22699 non-null  int64
 9   DOLocationID          22699 non-null  int64
10   payment_type          22699 non-null  int64
11   fare_amount           22699 non-null  float64
12   extra                 22699 non-null  float64
13   mta_tax               22699 non-null  float64
14   tip_amount            22699 non-null  float64
15   tolls_amount          22699 non-null  float64
16   improvement_surcharge 22699 non-null  float64
17   total_amount          22699 non-null  float64
dtypes: float64(8), int64(7), object(3)
memory usage: 3.1+ MB
None
```

Use `.head()`.

```
[92]: df.head()
```

```
[92]: Unnamed: 0  VendorID      tpep_pickup_datetime  tpep_dropoff_datetime  \
0      24870114          2  03/25/2017 8:55:43 AM  03/25/2017 9:09:47 AM
1      35634249          1  04/11/2017 2:53:28 PM  04/11/2017 3:19:58 PM
2     106203690          1  12/15/2017 7:26:56 AM  12/15/2017 7:34:08 AM
3      38942136          2  05/07/2017 1:17:59 PM  05/07/2017 1:48:14 PM
4      30841670          2  04/15/2017 11:32:20 PM  04/15/2017 11:49:03 PM

      passenger_count  trip_distance  RatecodeID  store_and_fwd_flag  \
0                   6           3.34           1                   N
1                   1           1.80           1                   N
2                   1           1.00           1                   N
3                   1           3.70           1                   N
4                   1           4.37           1                   N

      PULocationID  DOLocationID  payment_type  fare_amount  extra  mta_tax  \
0              100           231           1          13.0    0.0    0.5
1              186           43           1          16.0    0.0    0.5
2              262           236           1           6.5    0.0    0.5
3              188           97           1          20.5    0.0    0.5
4               4           112           2          16.5    0.5    0.5

      tip_amount  tolls_amount  improvement_surcharge  total_amount
0           2.76           0.0              0.3          16.56
1           4.00           0.0              0.3          20.80
2           1.45           0.0              0.3           8.75
3           6.39           0.0              0.3          27.69
4           0.00           0.0              0.3          17.80
```

Create trip_duration.

```
[93]: df["drop_off_converted"] = pd.to_datetime(df["tpep_dropoff_datetime"],
        ↪format="%m/%d/%Y %I:%M:%S %p")

df["pick_up_converted"] = pd.to_datetime(df["tpep_pickup_datetime"], format="%m/
        ↪%d/%Y %I:%M:%S %p")

df['trip_duration']=(df['drop_off_converted']-df['pick_up_converted'])/np.
        ↪timedelta64(1,"m")
df.head(10)
```

```
[93]: Unnamed: 0  VendorID      tpep_pickup_datetime  tpep_dropoff_datetime  \
0      24870114          2  03/25/2017 8:55:43 AM  03/25/2017 9:09:47 AM
1      35634249          1  04/11/2017 2:53:28 PM  04/11/2017 3:19:58 PM
2     106203690          1  12/15/2017 7:26:56 AM  12/15/2017 7:34:08 AM
3      38942136          2  05/07/2017 1:17:59 PM  05/07/2017 1:48:14 PM
4      30841670          2  04/15/2017 11:32:20 PM  04/15/2017 11:49:03 PM
5      23345809          2  03/25/2017 8:34:11 PM  03/25/2017 8:42:11 PM
```

6	37660487	2	05/03/2017 7:04:09 PM	05/03/2017 8:03:47 PM
7	69059411	2	08/15/2017 5:41:06 PM	08/15/2017 6:03:05 PM
8	8433159	2	02/04/2017 4:17:07 PM	02/04/2017 4:29:14 PM
9	95294817	1	11/10/2017 3:20:29 PM	11/10/2017 3:40:55 PM

	passenger_count	trip_distance	RatecodeID	store_and_fwd_flag	\
0	6	3.34	1	N	
1	1	1.80	1	N	
2	1	1.00	1	N	
3	1	3.70	1	N	
4	1	4.37	1	N	
5	6	2.30	1	N	
6	1	12.83	1	N	
7	1	2.98	1	N	
8	1	1.20	1	N	
9	1	1.60	1	N	

	PULocationID	DOLocationID	...	fare_amount	extra	mta_tax	tip_amount	\
0	100	231	...	13.0	0.0	0.5	2.76	
1	186	43	...	16.0	0.0	0.5	4.00	
2	262	236	...	6.5	0.0	0.5	1.45	
3	188	97	...	20.5	0.0	0.5	6.39	
4	4	112	...	16.5	0.5	0.5	0.00	
5	161	236	...	9.0	0.5	0.5	2.06	
6	79	241	...	47.5	1.0	0.5	9.86	
7	237	114	...	16.0	1.0	0.5	1.78	
8	234	249	...	9.0	0.0	0.5	0.00	
9	239	237	...	13.0	0.0	0.5	2.75	

	tolls_amount	improvement_surcharge	total_amount	drop_off_converted	\
0	0.0	0.3	16.56	2017-03-25 09:09:47	
1	0.0	0.3	20.80	2017-04-11 15:19:58	
2	0.0	0.3	8.75	2017-12-15 07:34:08	
3	0.0	0.3	27.69	2017-05-07 13:48:14	
4	0.0	0.3	17.80	2017-04-15 23:49:03	
5	0.0	0.3	12.36	2017-03-25 20:42:11	
6	0.0	0.3	59.16	2017-05-03 20:03:47	
7	0.0	0.3	19.58	2017-08-15 18:03:05	
8	0.0	0.3	9.80	2017-02-04 16:29:14	
9	0.0	0.3	16.55	2017-11-10 15:40:55	

	pick_up_converted	trip_duration
0	2017-03-25 08:55:43	14.066667
1	2017-04-11 14:53:28	26.500000
2	2017-12-15 07:26:56	7.200000
3	2017-05-07 13:17:59	30.250000
4	2017-04-15 23:32:20	16.716667

```

5 2017-03-25 20:34:11      8.000000
6 2017-05-03 19:04:09     59.633333
7 2017-08-15 17:41:06     21.983333
8 2017-02-04 16:17:07     12.116667
9 2017-11-10 15:20:29     20.433333

```

[10 rows x 21 columns]

Check for missing data and duplicates using `.isna()` and `.drop_duplicates()`.

```

[94]: # Check for missing data and duplicates using .isna() and .drop_duplicates()
print('Shape before removing duplicates',df.shape)
print('Shape after removing duplicates',df.drop_duplicates().shape)
df.isna().sum().sum()

```

Shape before removing duplicates (22699, 21)

Shape after removing duplicates (22699, 21)

[94]: 0

```

[95]: df.describe()

```

```

[95]:      Unnamed: 0      VendorID  passenger_count  trip_distance  \
count  2.269900e+04  22699.000000      22699.000000      22699.000000
mean    5.675849e+07      1.556236          1.642319          2.913313
std     3.274493e+07      0.496838          1.285231          3.653171
min     1.212700e+04      1.000000          0.000000          0.000000
25%     2.852056e+07      1.000000          1.000000          0.990000
50%     5.673150e+07      2.000000          1.000000          1.610000
75%     8.537452e+07      2.000000          2.000000          3.060000
max     1.134863e+08      2.000000          6.000000          33.960000

      RatecodeID  PULocationID  DOLocationID  payment_type  fare_amount  \
count  22699.000000  22699.000000  22699.000000  22699.000000  22699.000000
mean     1.043394    162.412353    161.527997     1.336887     13.026629
std     0.708391     66.633373     70.139691     0.496211     13.243791
min     1.000000     1.000000     1.000000     1.000000    -120.000000
25%     1.000000    114.000000    112.000000     1.000000      6.500000
50%     1.000000    162.000000    162.000000     1.000000      9.500000
75%     1.000000    233.000000    233.000000     2.000000     14.500000
max     99.000000    265.000000    265.000000     4.000000     999.990000

      extra      mta_tax      tip_amount  tolls_amount  \
count  22699.000000  22699.000000  22699.000000  22699.000000
mean     0.333275     0.497445     1.835781     0.312542
std     0.463097     0.039465     2.800626     1.399212
min    -1.000000    -0.500000     0.000000     0.000000


```

25%	0.000000	0.500000	0.000000	0.000000
50%	0.000000	0.500000	1.350000	0.000000
75%	0.500000	0.500000	2.450000	0.000000
max	4.500000	0.500000	200.000000	19.100000

	improvement_surcharge	total_amount	trip_duration
count	22699.000000	22699.000000	22699.000000
mean	0.299551	16.310502	17.013777
std	0.015673	16.097295	61.996482
min	-0.300000	-120.300000	-16.983333
25%	0.300000	8.750000	6.650000
50%	0.300000	11.800000	11.183333
75%	0.300000	17.800000	18.383333
max	0.300000	1200.290000	1439.550000

Create a scatterplot to visualize the relationship between variables of interest.


```
[96]: # Create a scatterplot to visualize the relationship between variables of
      ↪ interest
sns.set(style='whitegrid')
f = plt.figure()
f.set_figwidth(15)
f.set_figheight(15)
sns.regplot(x=df["trip_duration"], y=df["trip_distance"])
plt.ylim(0, 140)
plt.xlim(0,140)
plt.show()
```



Create a pairplot to visualize pairwise relationships between relevant variables.

```
[97]: # Create a pairplot to visualize pairwise relationships between variables in
      ↪ the data
      sns.pairplot(df[['trip_distance', 'tip_amount', 'total_amount',
      ↪ 'trip_duration']])
```

```
[97]: <seaborn.axisgrid.PairGrid at 0x7f51fa0e5c10>
```

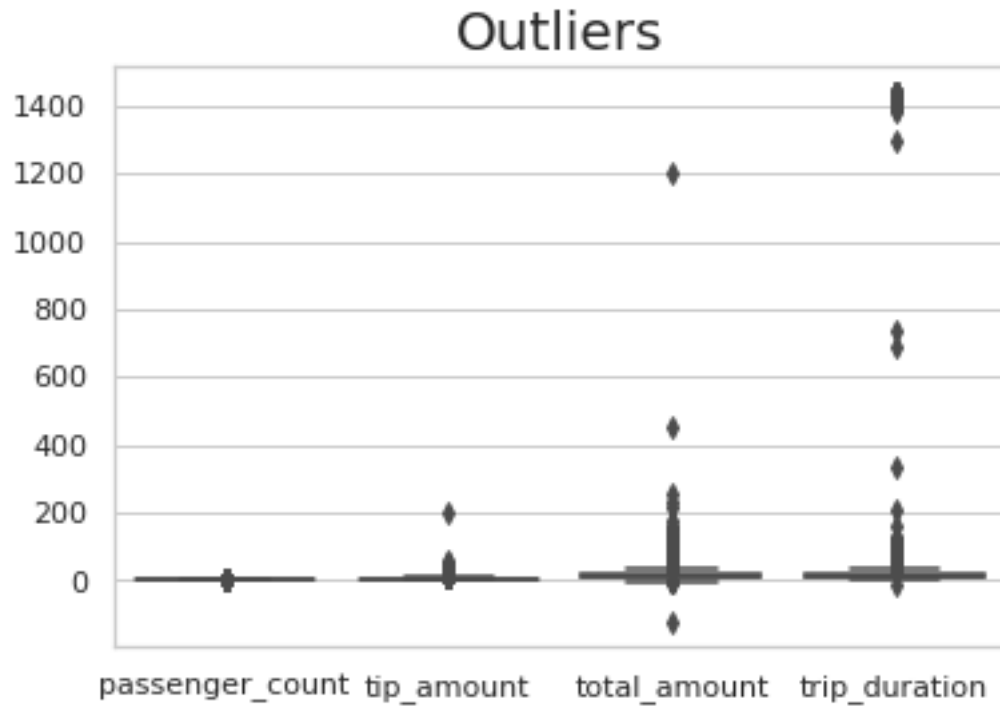



6.2 Task 2b. Address any outliers

Use a boxplot to visualize any outliers.

```
[98]: # Create boxplot to visualize the outliers
g = sns.boxplot(data=df[["passenger_count", "tip_amount", "total_amount",
    ↪ "trip_duration"]], showfliers=True);
g.set_title("Outliers", fontsize=20)
```

```
[98]: Text(0.5, 1.0, 'Outliers')
```

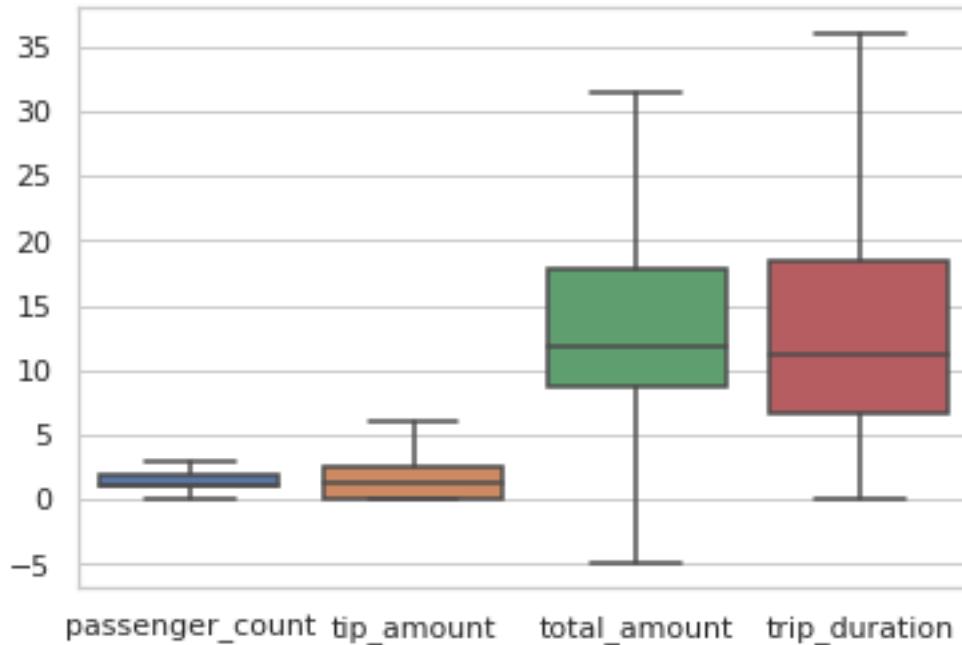


Use a boxplot to visualize the distribution of the data without outliers.

```
[99]: # Create boxplot to visualize distribution of data without outliers
g = sns.boxplot(data=df[["passenger_count", "tip_amount", "total_amount", "trip_duration"]], showfliers=False);
g.set_title("4 Variables without Outliers", fontsize=20)
```

```
[99]: Text(0.5, 1.0, '4 Variables without Outliers')
```

4 Variables without Outliers



Remove outliers as needed.

```
[100]: dpercentile_25 = df['trip_duration'].quantile(0.25)
dpercentile_75 = df['trip_duration'].quantile(0.75)
iqr= dpercentile_75 - dpercentile_25
upper_limit = dpercentile_75 + 1.5 * iqr

df[df["trip_duration"] > upper_limit] = upper_limit
df[df["trip_duration"] < 0] = 0

apercentile_25 = df["total_amount"].quantile(0.25)
apercentile_75 = df["total_amount"].quantile(0.75)
iqr= apercentile_75 - apercentile_25
aupper_limit = apercentile_75 + 1.5 * iqr

df[df["total_amount"] > aupper_limit] = aupper_limit
df[df["total_amount"] < 0] = 0
```

6.3 Task 2c. Identify correlations

Next, code a correlation matrix to help determine most correlated variables.

```
[101]: df.corr(method='pearson')
```

[101]:

	Unnamed: 0	VendorID	passenger_count	trip_distance	\
Unnamed: 0	1.000000	-0.484018	-0.480060	-0.479913	
VendorID	-0.484018	1.000000	0.991630	0.986383	
passenger_count	-0.480060	0.991630	1.000000	0.979200	
trip_distance	-0.479913	0.986383	0.979200	1.000000	
RatecodeID	-0.484752	0.998592	0.991031	0.987447	
PULocationID	0.249919	-0.531011	-0.526598	-0.533220	
DOLocationID	0.254780	-0.517895	-0.512815	-0.524301	
payment_type	-0.483891	0.997301	0.989897	0.986066	
fare_amount	-0.388750	0.812167	0.807055	0.885940	
extra	-0.484724	0.997964	0.990369	0.987087	
mta_tax	-0.484852	0.998656	0.991107	0.987586	
tip_amount	-0.480481	0.988855	0.981257	0.985027	
tolls_amount	-0.484900	0.998252	0.990722	0.987633	
improvement_surcharge	-0.484874	0.998655	0.991107	0.987580	
total_amount	-0.350981	0.733121	0.728428	0.818807	
trip_duration	-0.309621	0.670141	0.666417	0.744132	

	RatecodeID	PULocationID	DOLocationID	payment_type	\
Unnamed: 0	-0.484752	0.249919	0.254780	-0.483891	
VendorID	0.998592	-0.531011	-0.517895	0.997301	
passenger_count	0.991031	-0.526598	-0.512815	0.989897	
trip_distance	0.987447	-0.533220	-0.524301	0.986066	
RatecodeID	1.000000	-0.530776	-0.517753	0.998632	
PULocationID	-0.530776	1.000000	0.349153	-0.530324	
DOLocationID	-0.517753	0.349153	1.000000	-0.517543	
payment_type	0.998632	-0.530324	-0.517543	1.000000	
fare_amount	0.813416	-0.459258	-0.459545	0.811169	
extra	0.999219	-0.530773	-0.518097	0.997943	
mta_tax	0.999908	-0.530843	-0.517718	0.998685	
tip_amount	0.990166	-0.526370	-0.514729	0.983980	
tolls_amount	0.999550	-0.530851	-0.518053	0.998322	
improvement_surcharge	0.999921	-0.530875	-0.517755	0.998688	
total_amount	0.734320	-0.419093	-0.423095	0.726262	
trip_duration	0.670954	-0.379278	-0.380769	0.668878	

	fare_amount	extra	mta_tax	tip_amount	\
Unnamed: 0	-0.388750	-0.484724	-0.484852	-0.480481	
VendorID	0.812167	0.997964	0.998656	0.988855	
passenger_count	0.807055	0.990369	0.991107	0.981257	
trip_distance	0.885940	0.987087	0.987586	0.985027	
RatecodeID	0.813416	0.999219	0.999908	0.990166	
PULocationID	-0.459258	-0.530773	-0.530843	-0.526370	
DOLocationID	-0.459545	-0.518097	-0.517718	-0.514729	
payment_type	0.811169	0.997943	0.998685	0.983980	
fare_amount	1.000000	0.812888	0.813251	0.835297	
extra	0.812888	1.000000	0.999308	0.989821	

mta_tax	0.813251	0.999308	1.000000	0.990257
tip_amount	0.835297	0.989821	0.990257	1.000000
tolls_amount	0.814062	0.998920	0.999605	0.989955
improvement_surcharge	0.813258	0.999309	0.999999	0.990254
total_amount	0.980331	0.735808	0.734164	0.780913
trip_duration	0.940505	0.670915	0.671400	0.702192


	tolls_amount	improvement_surcharge	total_amount	\
Unnamed: 0	-0.484900	-0.484874	-0.350981	
VendorID	0.998252	0.998655	0.733121	
passenger_count	0.990722	0.991107	0.728428	
trip_distance	0.987633	0.987580	0.818807	
RatecodeID	0.999550	0.999921	0.734320	
PULocationID	-0.530851	-0.530875	-0.419093	
DOLocationID	-0.518053	-0.517755	-0.423095	
payment_type	0.998322	0.998688	0.726262	
fare_amount	0.814062	0.813258	0.980331	
extra	0.998920	0.999309	0.735808	
mta_tax	0.999605	0.999999	0.734164	
tip_amount	0.989955	0.990254	0.780913	
tolls_amount	1.000000	0.999609	0.736071	
improvement_surcharge	0.999609	1.000000	0.734161	
total_amount	0.736071	0.734161	1.000000	
trip_duration	0.671988	0.671348	0.941669	

	trip_duration
Unnamed: 0	-0.309621
VendorID	0.670141
passenger_count	0.666417
trip_distance	0.744132
RatecodeID	0.670954
PULocationID	-0.379278
DOLocationID	-0.380769
payment_type	0.668878
fare_amount	0.940505
extra	0.670915
mta_tax	0.671400
tip_amount	0.702192
tolls_amount	0.671988
improvement_surcharge	0.671348
total_amount	0.941669
trip_duration	1.000000

Visualize a correlation heatmap of the data.

```
[102]: # Create correlation heatmap
plt.figure(figsize=(15,15))
```

```
sns.heatmap(df.corr(method="pearson"), annot=True, cmap="Blues")
plt.title("Correlation heatmap indicates many low correlated variables",
          fontsize=18)
plt.show()
```



7 PACE: Construct

After analysis and deriving variables with close relationships, it is time to begin constructing the model. Consider the questions in the [PACE Strategy Doc](#) to reflect on the Constructing stage of this task. * Why did you select the X variables you did?

Dropped columns based on correlations between variables and multicollinearity, and fine-tuned by running and rerunning models to examine change in R^2 , MAE, and RMSE.

7.1 Task 3a. Select outcome variable and features

Set your Y and X variables. Y represents the outcome variable, and X represents the features.

```
[103]: Y = df[["trip_duration"]]
X = df.drop(columns="trip_duration")
X.head()
```

```
[103]:   Unnamed: 0  VendorID  tpep_pickup_datetime  tpep_dropoff_datetime \
0   24870114.0        2.0  03/25/2017 8:55:43 AM  03/25/2017 9:09:47 AM
1   35634249.0        1.0  04/11/2017 2:53:28 PM  04/11/2017 3:19:58 PM
2  106203690.0        1.0  12/15/2017 7:26:56 AM  12/15/2017 7:34:08 AM
3   38942136.0        2.0  05/07/2017 1:17:59 PM  05/07/2017 1:48:14 PM
4   30841670.0        2.0  04/15/2017 11:32:20 PM  04/15/2017 11:49:03 PM
```

```
   passenger_count  trip_distance  RatecodeID  store_and_fwd_flag \
0                6.0           3.34         1.0                 N
1                1.0           1.80         1.0                 N
2                1.0           1.00         1.0                 N
3                1.0           3.70         1.0                 N
4                1.0           4.37         1.0                 N
```

```
   PULocationID  DOLocationID  payment_type  fare_amount  extra  mta_tax \
0           100.0          231.0         1.0         13.0    0.0    0.5
1           186.0           43.0         1.0         16.0    0.0    0.5
2           262.0          236.0         1.0          6.5    0.0    0.5
3           188.0           97.0         1.0         20.5    0.0    0.5
4            4.0          112.0         2.0         16.5    0.5    0.5
```

```
   tip_amount  tolls_amount  improvement_surcharge  total_amount \
0         2.76          0.0                 0.3         16.56
1         4.00          0.0                 0.3         20.80
2         1.45          0.0                 0.3          8.75
3         6.39          0.0                 0.3         27.69
4         0.00          0.0                 0.3         17.80
```

```
   drop_off_converted  pick_up_converted
0  2017-03-25 09:09:47  2017-03-25 08:55:43
1  2017-04-11 15:19:58  2017-04-11 14:53:28
2  2017-12-15 07:34:08  2017-12-15 07:26:56
3  2017-05-07 13:48:14  2017-05-07 13:17:59
4  2017-04-15 23:49:03  2017-04-15 23:32:20
```

7.2 Task 3b. Pre-process data

To help with processing time, consider dropping irrelevant and redundant columns.

```
[104]: columns_to_drop = ['tpep_pickup_datetime', 'tpep_dropoff_datetime',
                        'store_and_fwd_flag', 'passenger_count', 'VendorID',
                        'fare_amount', 'PULocationID', 'DOLocationID',
                        'total_amount',
                        'drop_off_converted', 'pick_up_converted']
X = X.drop(columns_to_drop, axis=1)
X = X.loc[:, ~X.columns.str.contains("Unnamed")]
X.head()
```

```
[104]:
```

	trip_distance	RatecodeID	payment_type	extra	mta_tax	tip_amount	\
0	3.34	1.0	1.0	0.0	0.5	2.76	
1	1.80	1.0	1.0	0.0	0.5	4.00	
2	1.00	1.0	1.0	0.0	0.5	1.45	
3	3.70	1.0	1.0	0.0	0.5	6.39	
4	4.37	1.0	2.0	0.5	0.5	0.00	

	tolls_amount	improvement_surcharge
0	0.0	0.3
1	0.0	0.3
2	0.0	0.3
3	0.0	0.3
4	0.0	0.3

Use `StandardScaler()` and `fit_transform()` to standardize the X variables.

```
[105]: scaler = StandardScaler()
x_scaled = scaler.fit_transform(X)
print("X scaled:", X_scaled)
```

```
X scaled: [[-0.18116118 -0.33999272 -0.37659905 ... -0.18555461 -0.34096692
 -0.33964103]
 [-0.35082338 -0.33999272 -0.37659905 ... -0.05118079 -0.34096692
 -0.33964103]
 [-0.43895959 -0.33999272 -0.37659905 ... -0.32751404 -0.34096692
 -0.33964103]
 ...
 [-0.50285834 -0.33999272 -0.26740498 ... -0.48464472 -0.34096692
 -0.33964103]
 [-0.28912804 -0.33999272 -0.37659905 ... -0.30042255 -0.34096692
 -0.33964103]
 [-0.31777231 -0.33999272 -0.37659905 ... -0.22998466 -0.34096692
 -0.33964103]]
```


7.3 Task 3c. Build model

Create training and testing sets.

```
[106]: x_train, x_test, y_train, y_test = train_test_split(x_scaled, Y, test_size=0.2, random_state=0)
```

Build and fit your model to the data.

```
[107]: lr.fit(x_train, y_train)
```

```
[107]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

7.4 Task 3d. Evaluate model

Evaluate your model performance by calculating the residual sum of squares and the explained variance score (R^2). Calculate the Mean Absolute Error, Mean Squared Error, and the Root Mean Squared Error.

```
[108]: r_sq = lr.score(x_train, y_train)
print("Coefficient of determination:", r_sq)
y_pred = lr.predict(x_train)
print("R^2:", r2_score(y_train, y_pred))
print("MAE:", mean_absolute_error(y_train, y_pred))
print("RMSE:", np.sqrt(mean_squared_error(y_train, y_pred)))
```

```
Coefficient of determination: 0.7391110948153115
R^2: 0.7391110948153116
MAE: 3.2119417836310444
RMSE: 4.5781331414091895
```

```
[109]: r_sq_test = lr.score(x_test, y_test)
print("Coefficient of determination:", r_sq_test)
y_pred_test = lr.predict(x_test)
print("R^2:", r2_score(y_test, y_pred_test))
print("MAE:", mean_absolute_error(y_test, y_pred_test))
print("RMSE:", np.sqrt(mean_squared_error(y_test, y_pred_test)))
```

```
Coefficient of determination: 0.7332003491146651
R^2: 0.7332003491146653
MAE: 3.2214448090342804
RMSE: 4.601147377111918
```

8 PACE: Execute

Consider these questions [PACE Strategy Doc](#) to reflect on the Execute stage of this task.

8.1 Task 4a. Results

If the linear regression assumptions are met, the model results can be appropriately interpreted.

Use the code cell below to get `actual`, `predicted`, and `residual` for the testing set, and store them as columns in a `results` dataframe.

```
[110]: results = pd.DataFrame(data={"actual": y_test["trip_duration"],
                                   "predicted": y_pred_test.ravel()})
results["residual"] = results["actual"] - results["predicted"]
results.head()
```


```
[110]:
```

	actual	predicted	residual
5818	18.016667	15.666455	2.350212
18134	31.375000	31.375124	-0.000124
4655	5.883333	7.334027	-1.450693
7378	15.950000	18.339652	-2.389652
13914	11.900000	11.543187	0.356813

8.2 Task 4b. Visualize model results


Create a scatterplot to visualize predicted over actual.

```
[111]: sns.set(style='whitegrid')
f = plt.figure()
f.set_figwidth(15)
f.set_figheight(15)
sns.regplot(x="actual",
            y="predicted",
            data=results, line_kws={"color": "red"})
plt.ylim(0, 40)
plt.xlim(0,40)
plt.show()
```




Visualize the distribution of the residuals.

```
[112]: # Visualize the distribution of the `residuals`  
      ### YOUR CODE HERE ###  
  
      plt.hist(results["residual"], bins=30)  
      plt.title("Distribution of the residuals")  
      plt.xlabel("residual value")  
      plt.ylabel("count")  
      plt.show()
```



Create a scatterplot of residuals over predicted.

```
[113]: sns.scatterplot(x="predicted", y="residual", data=results)
plt.axhline(0)
plt.title("Scatterplot of residuals over predicted values")
plt.xlabel("predicted value")
plt.ylabel("residual value")
plt.show()
```



8.3 Task 4c. Conclusion

1. What results can be presented from this notebook?
1. It would be important to present the model's assumptions and precision of the model through the MAE and the RMSE